

6.54 AAD

L10 Red Black Trees

This lecture: PFDs Ch 3.3

Red Black Trees


The binary search trees we studied in L4 may be highly unbalanced.

Red Black trees are binary search trees that remain approximately balanced.

Def: A Red Black tree is a binary search tree where each node is coloured either red or black such that

- (1) No red node has a red child.
- (2) Every path from the root to the leaf contains the same number of black nodes.

Ex :



How balanced are Red Black Trees?

Lemma

The maximum depth of a node in a Red Black Tree of size n is no more than $2 \lfloor \log(n+1) \rfloor$.

How can we prove this?

1) What information do we have?

Let s be the length of the shortest path,
 l be the length of the longest path,
 k be the number of black nodes along any path.

By (2) $k \leq s \leq l \leq 2k$, so $2s \geq l$



How large is l ?

Since each path contains at least s nodes,
the number of nodes in the tree is $n \geq 2^{s+1} - 1$.

Hence

$$2 \lfloor \log(n+1) \rfloor \geq 2 \lfloor \log(2^{s+1}) \rfloor$$

$$= 2 \lfloor s+1 \rfloor$$

$$\geq 2 \lfloor l/2 + 1 \rfloor$$

$$\geq l.$$

□

Operations on Red Black Trees

We assume empty nodes to be black.

data Colour = R | B

data RedBlackTree a = L | N Colour a

(RedBlackTree a) -- left
(RedBlackTree a) -- right

- 1) member :: a → RedBlackTree a → RedBlackTree a
member x ts ≡ does ts contain a node with value x?

member $O(\log n)$ by Lemma 1.
implemented as for binary search trees.

- 2) insert :: a → RedBlackTree a → RedBlackTree a
How do we manage to satisfy the two invariants?

- insert red leaf node

⇒ Number of black nodes unchanged,
so invariant 2 still satisfied.

⇒ Invariant 1 may be violated
in exactly one place in the tree.

ins x E = RB R x E E

ins x (RB c y left right)

$1 \leq y$ = balance (RB c y (ins x left) right)

$1 < y$ = balance (RB c x left (ins y right))

Makes sure
that root is
black.


{ insert Tree x ts =
RB B y left right
where RB - y left right = ins x ts.

Assume that invariance 1 is violated.


What will the tree look like, and how can we rebalance it?

CASE 1:

BEFORE BALANCING



BALANCE
REWRITING
STEP




- Invariant 2 is not violated, so the number of black nodes along the paths from a, b, c, d must be the same, say k .
- By the binary search tree property, we have

$$a \leq x \leq b \leq y \leq c \leq z \leq d$$


↑
Rebalance
along this node!

- After the balance rewriting step, invariant 1 may be violated at one higher level in the tree.
- Invariant 2 is satisfied.

CASE 2




BALANCE
REWRITING
STEP




$a \leq x \leq b \leq y \leq c \leq z \leq d$

↑
Rebalance
along this
node!

CASE 3




BALANCE
REWRITING
STEP
→




$a \leq x \leq b \leq y \leq c \leq z \leq d$

↑
Rebalance
along this
node!

CASE 4



BALANCE
REWRITING
STEP \rightarrow



$a \leq x \leq b \leq \gamma \leq c \leq z \leq d$

\uparrow
Rebalance
along this
node!